
Direct Self-evolving Optimization: Efficient Self-Evolving LLM Training without Challenger Model

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 The self-evolving training paradigm has emerged as a crucial technique for en-
2 hancing the ability of Large Language Models (LLMs) to perform mathematical
3 reasoning and handle complex tasks. By training a challenger model to generate
4 questions for a solver model, this approach effectively overcomes the scarcity of
5 human-annotated data. However, a major drawback of this paradigm is the signifi-
6 cant increase in computational and memory costs. In this paper, we eliminate the
7 need for challenger model training by deriving a closed-form optimal challenger
8 policy, conditioned on the current solver model. We show that this optimal policy
9 is a reweighted version of the base model’s policy, where the weights correspond
10 to the challenger’s reward. Finally, we employ the Metropolis-Hastings algorithm
11 to iteratively sample questions from this optimal policy, which are subsequently
12 fed into the solver model for efficient training.

13 1 Introduction

14 1.1 Existing method

15 In the existing self-evolving LLM training [Huang et al., 2025, Xia et al., 2025], the **challenger**
16 optimizes:

$$\max_{\phi} R_c(\phi, \theta) := \mathbb{E}_{x \sim \pi_{\phi}} [r_c(x, \theta)] - \beta \text{KL}(\pi_{\phi} || \pi_{ref}) \quad (1)$$

17 where $r_c(x, \theta)$ is the challenger reward function given a solver model, e.g., if a generated question x
18 makes solver’s answers more uncertain, challenger will get higher reward.

19 The **solver** optimizes:

$$\max_{\theta} R_s(\phi, \theta) := \mathbb{E}_{x \sim \pi_{\phi}} [\mathbb{E}_{y \sim \pi_{\theta}(y|x)} [r_s(x, y)]] \quad (2)$$

20 where $r_s(x, y)$ is the solver’s reward function.

21 Since the two objectives depend on both ϕ and θ , existing solution is to fix one and update another.
22 Given initialization ϕ_0 and θ_0 , the updating rule is

$$\begin{aligned} \phi_{t+1} &\leftarrow \text{update } \phi_t \text{ for K steps using GRPO on reward } R_c(\phi_t, \theta_t), \\ \theta_{t+1} &\leftarrow \text{update } \theta_t \text{ for K steps using GRPO on reward } R_s(\phi_{t+1}, \theta_t) \end{aligned}$$

23 1.2 The Weakness and Our proposal

24 However, training two models are costly and this alternating optimization is unstable. To optimize
25 equation 1 and equation 2 is to first define

$$\phi^*(\theta) := \arg \max_{\phi} R_c(\phi, \theta),$$

26 Indeed the closed form solution is

$$\pi_{\phi^*}(\theta) = \frac{1}{Z(\theta)} \pi_{ref}(x) \exp \frac{r_c(x, \theta)}{\beta}$$

27 then plug $\pi_{\phi^*}(\theta)$ into equation 2:

$$\begin{aligned} \max_{\theta} R(\theta) &:= R_s(\phi^*(\theta), \theta) \\ &= \mathbb{E}_{x \sim \pi_{\phi}} [\mathbb{E}_{y \sim \pi_{\theta}} [r_s(x, y)]] = \sum_x \pi_{\phi^*}(\theta)(x|p_0) \mathbb{E}_{x \sim \pi_{\theta}} [r_s(x, y)] \\ &= \sum_x \frac{1}{Z(\theta)} \pi_{ref}(x|p_0) \exp \frac{r_c(x, \theta)}{\beta} \mathbb{E}_{y \sim \pi_{\theta}(y|x)} [r_s(x, y)] \\ &= \frac{1}{Z(\theta)} \mathbb{E}_{x \sim \pi_{ref}(x|p_0)} \left[\exp \frac{r_c(x, \theta)}{\beta} \mathbb{E}_{y \sim \pi_{\theta}(y|x)} [r_s(x, y)] \right] \end{aligned} \quad (3)$$

28 where $r_c(x, \theta) = \max 0, r_{\text{uncertainty}}(x; \theta) - r_{\text{rep}}(x)$, $r_{\text{uncertainty}}(x; \theta) := 1 - 2|\hat{p}(x; S_{\theta}(x)) - \frac{1}{2}|$, and
29 $\hat{p}(x; S_{\theta}(x)) = \frac{1}{m} \sum_{j=1}^m \mathbb{I} y_j = \bar{y}(x)$, and $S_{\theta}(x) = y_1, \dots, y_m$ is the set of m responses that solver
30 model θ generated for question x , and $\bar{y}(x)$ is the most frequent answer in $S_{\theta}(x)$.

31 $r_{\text{rep}}(x)$ is the repetition penalty. Given a batch of questions $x_1, \dots, x_B$, we compute pairwise
32 distances using BLEU score similarity, $d_{ij} = 1 - BLEU(x_i, x_j)$, and group questions where
33 $d_{ij} < \tau_{BLEU}$ into clusters $C = \{C_1, \dots, C_K\}$. The penalty for a question x_i in a cluster C_k is
34 proportional to its relative size:

$$r_{\text{rep}}(x) = \frac{|C_k|}{B}.$$

35 For $r_s(x, y_j)$, it is 1 if y_j is identical to the pseudo label $\bar{y}(x)$, else 0.

36 Hence, we can directly optimize this objective using policy gradient. First, we generate a mini-batch
37 of questions from the base model $\pi_{ref}: \{x_1, \dots, x_G\}$. For each question, the solver model generates
38 m responses: $A = \{y_{i,1}, \dots, y_{i,m}\}$. Then we estimate the policy gradient as

$$\hat{\nabla}_{\theta} R(\theta) = \frac{1}{G} \sum_{i=1}^G \exp \frac{r_c(x_i, \theta)}{\beta} \frac{1}{m} \sum_{j=1}^m \nabla_{\theta} \log \pi_{\theta}(y_{i,j}|x_i) \quad (4)$$

39 We directly optimize the solver model. It is a weighted policy gradient method, where the weights
40 depend on how the current solver model performs on the generated questions. For the questions that
41 the solver has higher uncertainty, its gradient will be up-weighted.

42 1.3 Active Sampling

43 The objective in equation 3 can be equivalently written as:

$$\max_{\theta} R(\theta) = \mathbb{E}_{x \sim \pi_{\phi^*}(\theta)(x|p_0)} [\mathbb{E}_{y \sim \pi_{\theta}(y|x)} [r_s(x, y)]]$$

44 where $\pi_{\phi^*}(\theta) \sim \pi_{ref}(x) \exp \frac{r_c(x, \theta)}{\beta}$. Now the question is, how can we actively sample questions
45 from $\pi_{\phi^*}(\theta)(x|p_0)$ only given access to π_{base} ? A generally used approach is MCMC sampling via
46 Metropolis-Hasting algorithm.

Algorithm 1: MCMC-DEO

Input: θ_0 , question-generating prompt p_0 , learning rate η , regularization parameter β .

for $t = 0, \dots, T - 1$ **do**

 Sample N questions $x_1^0, \dots, x_N^0$ from $\pi_{base}(\cdot|p_0)$.

 For each x_i^0 , generate mutated version $\hat{x}_i = MCMC(x_i^0, \tau, \pi_{base})$ For each $\hat{x}_i$, generate m answers $y_{i,1}, \dots, y_{i,m}$ from $\pi_{\theta_t}(\cdot|\hat{x}_i)$.

 Compute mini-batch gradient $\hat{\nabla}_{\theta} R(\theta) = \frac{1}{N} \sum_{i=1}^N \frac{1}{m} \sum_{j=1}^m \nabla_{\theta} r_s(x_i, y_{i,j}) \log \pi_{\theta}(y_{i,j}|\hat{x}_i)$.

$\theta_{t+1} = \theta_t + \eta \hat{\nabla}_{\theta} R(\theta)$.

end

Output: θ_T

Algorithm 2: MCMC(x_0, τ, π)

Input: Original question x_0 , mutation prompt p , policy π

for $t = 0, \dots, \tau - 1$ **do**

 Sample a mutated question $x_{t+1} \sim \pi(x_t|p)$,

 Compute acceptance probability $\alpha \doteq \min 1, \exp(r_c(x') - r_c(x))$

 Accepting $x_{t+1} = x'$ with probability α , otherwise $x_{t+1} = x_t$.

end

Output: x_{τ}

References

- 47 Chengsong Huang, Wenhao Yu, Xiaoyang Wang, Hongming Zhang, Zongxia Li, Ruosen Li, Jiaxin
48 Huang, Haitao Mi, and Dong Yu. R-zero: Self-evolving reasoning llm from zero data. *arXiv*
49 *preprint arXiv:2508.05004*, 2025.
- 51 Peng Xia, Kaide Zeng, Jiaqi Liu, Can Qin, Fang Wu, Yiyang Zhou, Caiming Xiong, and Huaxiu
52 Yao. Agent0: Unleashing self-evolving agents from zero data via tool-integrated reasoning. *arXiv*
53 *preprint arXiv:2511.16043*, 2025.

54 NeurIPS Paper Checklist

55 The checklist is designed to encourage best practices for responsible machine learning research,
56 addressing issues of reproducibility, transparency, research ethics, and societal impact. Do not remove
57 the checklist: **The papers not including the checklist will be desk rejected.** The checklist should
58 follow the references and follow the (optional) supplemental material. The checklist does NOT count
59 towards the page limit.

60 Please read the checklist guidelines carefully for information on how to answer these questions. For
61 each question in the checklist:

- 62 • You should answer [Yes], [No], or [N/A].
- 63 • [N/A] means either that the question is Not Applicable for that particular paper or the
64 relevant information is Not Available.
- 65 • Please provide a short (1–2 sentence) justification right after your answer (even for [N/A]).

66 **The checklist answers are an integral part of your paper submission.** They are visible to the
67 reviewers, area chairs, senior area chairs, and ethics reviewers. You will also be asked to include it
68 (after eventual revisions) with the final version of your paper, and its final version will be published
69 with the paper.

70 The reviewers of your paper will be asked to use the checklist as one of the factors in their evaluation.
71 While [Yes] is generally preferable to [No], it is perfectly acceptable to answer [No] provided a
72 proper justification is given (e.g., error bars are not reported because it would be too computationally
73 expensive” or “we were unable to find the license for the dataset we used”). In general, answering
74 [No] or [N/A] is not grounds for rejection. While the questions are phrased in a binary way, we
75 acknowledge that the true answer is often more nuanced, so please just use your best judgment and
76 write a justification to elaborate. All supporting evidence can appear either in the main paper or the
77 supplemental material, provided in appendix. If you answer [Yes] to a question, in the justification
78 please point to the section(s) where related material for the question can be found.

79 IMPORTANT, please:

- 80 • **Delete this instruction block, but keep the section heading “NeurIPS Paper Checklist”.**
- 81 • **Keep the checklist subsection headings, questions/answers and guidelines below.**
- 82 • **Do not modify the questions and only use the provided macros for your answers.**

83 1. Claims

84 Question: Do the main claims made in the abstract and introduction accurately reflect the
85 paper’s contributions and scope?

86 Answer: [TODO]

87 Justification: [TODO]

88 Guidelines:

- 89 • The answer [N/A] means that the abstract and introduction do not include the claims
90 made in the paper.
- 91 • The abstract and/or introduction should clearly state the claims made, including the
92 contributions made in the paper and important assumptions and limitations. A [No] or
93 [N/A] answer to this question will not be perceived well by the reviewers.
- 94 • The claims made should match theoretical and experimental results, and reflect how
95 much the results can be expected to generalize to other settings.
- 96 • It is fine to include aspirational goals as motivation as long as it is clear that these goals
97 are not attained by the paper.

98 2. Limitations

99 Question: Does the paper discuss the limitations of the work performed by the authors?

100 Answer: [TODO]

101 Justification: [TODO]

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that the paper does not include experiments.

- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).

- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer **[Yes]** if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.

- 258 • The paper should disclose whether the full research project required more compute
259 than the experiments reported in the paper (e.g., preliminary or failed experiments that
260 didn't make it into the paper).

261 **9. Code of ethics**

262 Question: Does the research conducted in the paper conform, in every respect, with the
263 NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines?>

264 Answer: **[TODO]**

265 Justification: **[TODO]**

266 Guidelines:

- 267 • The answer [N/A] means that the authors have not reviewed the NeurIPS Code of
268 Ethics.
- 269 • If the authors answer [No], they should explain the special circumstances that require a
270 deviation from the Code of Ethics.
- 271 • The authors should make sure to preserve anonymity (e.g., if there is a special consid-
272 eration due to laws or regulations in their jurisdiction).

273 **10. Broader impacts**

274 Question: Does the paper discuss both potential positive societal impacts and negative
275 societal impacts of the work performed?

276 Answer: **[TODO]**

277 Justification: **[TODO]**

278 Guidelines:

- 279 • The answer [N/A] means that there is no societal impact of the work performed.
- 280 • If the authors answer [N/A] or [No], they should explain why their work has no societal
281 impact or why the paper does not address societal impact.
- 282 • Examples of negative societal impacts include potential malicious or unintended uses
283 (e.g., disinformation, generating fake profiles, surveillance), fairness considerations
284 (e.g., deployment of technologies that could make decisions that unfairly impact specific
285 groups), privacy considerations, and security considerations.
- 286 • The conference expects that many papers will be foundational research and not tied
287 to particular applications, let alone deployments. However, if there is a direct path to
288 any negative applications, the authors should point it out. For example, it is legitimate
289 to point out that an improvement in the quality of generative models could be used to
290 generate Deepfakes for disinformation. On the other hand, it is not needed to point out
291 that a generic algorithm for optimizing neural networks could enable people to train
292 models that generate Deepfakes faster.
- 293 • The authors should consider possible harms that could arise when the technology is
294 being used as intended and functioning correctly, harms that could arise when the
295 technology is being used as intended but gives incorrect results, and harms following
296 from (intentional or unintentional) misuse of the technology.
- 297 • If there are negative societal impacts, the authors could also discuss possible mitigation
298 strategies (e.g., gated release of models, providing defenses in addition to attacks,
299 mechanisms for monitoring misuse, mechanisms to monitor how a system learns from
300 feedback over time, improving the efficiency and accessibility of ML).

301 **11. Safeguards**

302 Question: Does the paper describe safeguards that have been put in place for responsible
303 release of data or models that have a high risk for misuse (e.g., pre-trained language models,
304 image generators, or scraped datasets)?

305 Answer: **[TODO]**

306 Justification: **[TODO]**

307 Guidelines:

- 308 • The answer [N/A] means that the paper poses no such risks.

- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. Declaration of LLM usage

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.